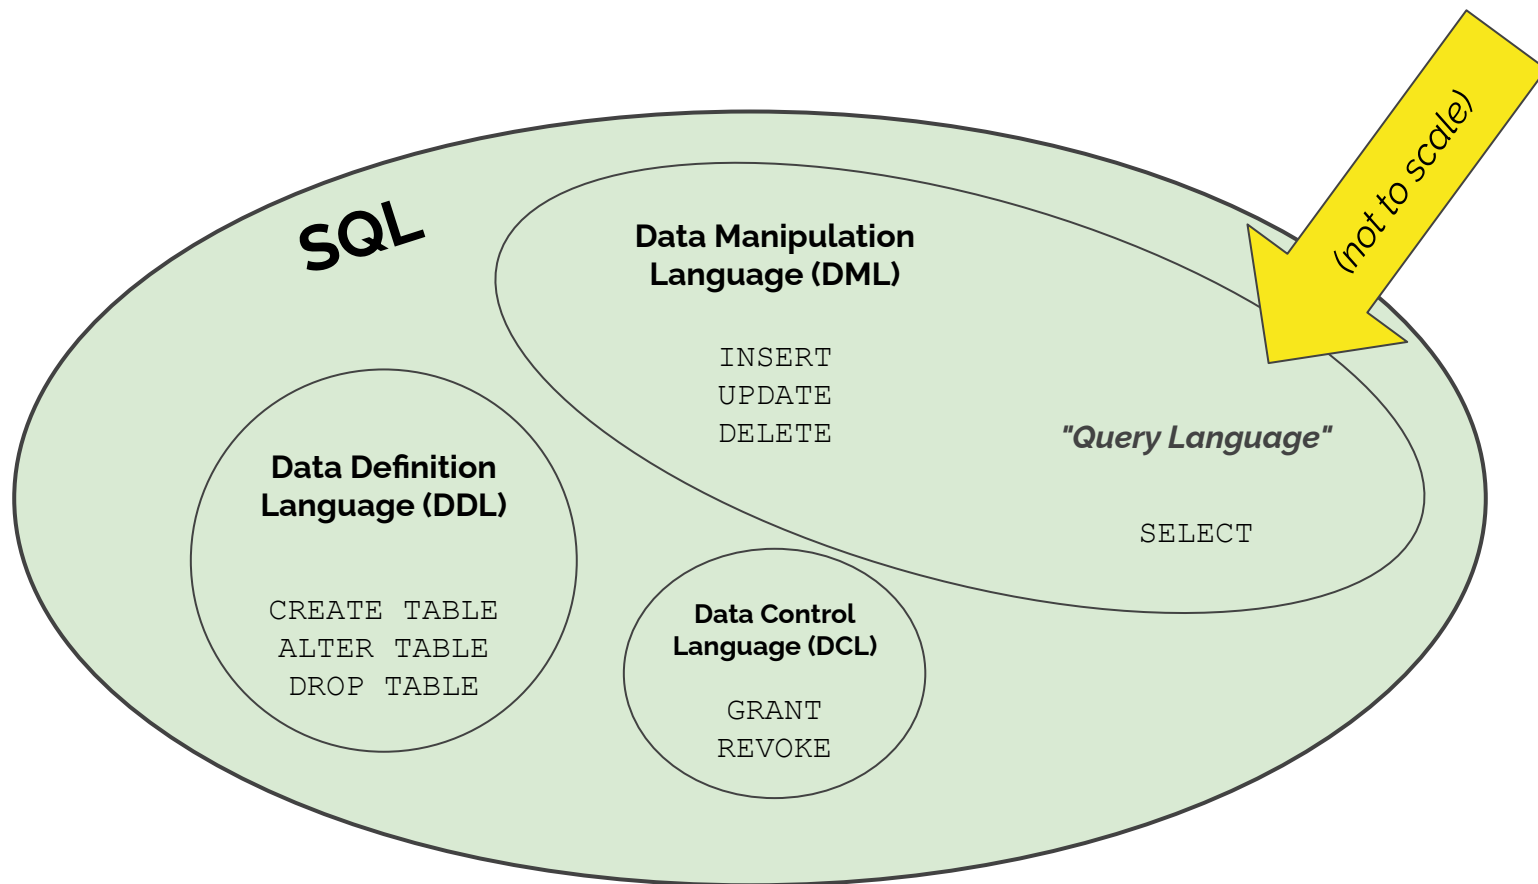


CSC 365

Introduction to Database Systems



- Five core relational algebra operators can be combined to express interesting and complex queries
 - Selection
 - Projection
 - Cartesian Product
 - Union
 - Difference
- A database query language (ie. SQL) should be *at least* as expressive as relational algebra
 - SQL-92: <http://www.contrib.andrew.cmu.edu/~shadow/sql/sql1992.txt>

Minimal syntax:

```
SELECT <list of columns>  
FROM <table(s)>  
[ WHERE <row filter expression> ]  
;
```

Asterisk (*) can be used in place of column list to indicate all columns.

WHERE clause is optional

AIRPLANE

<u>TailNum</u>	Make	Model	MaxSpeed
C97W	Boeing	797	<i>null</i>
R53Q	Cessna	FG	220
T80H	Airbus	A380	634
G59K	Airbus	A320	450
P88T	Piper	Arrow	180
K30W	Boeing	707	450

FLIGHT

<u>TailNum</u>	<u>PilotID</u> -----	<u>CopilotID</u> -----	Runway	<u>Date</u>
R53Q	K407	D342	S-2	9/1/17
T80H	K407	<i>null</i>	W-2	9/21/17
C97W	D342	<i>null</i>	W-2	8/9/21
T80H	D342	K407	W-3	9/9/17

PILOT

<u>PilotID</u>	Name
D342	Charlie
K407	Juliett
H452	Piper

```
SELECT Make, Model, MaxSpeed  
FROM Airplane;
```

Column name(s)

```
SELECT Make, Model, MaxSpeed  
FROM Airplane  
WHERE MaxSpeed > 500;
```

Table name(s)

Row selection
expression

Relational Algebra Operator		SQL
π	Projection	SELECT DISTINCT
\times	Cartesian product	<i>comma</i> or CROSS JOIN
\bowtie_{θ}	Theta join	INNER JOIN ... ON
σ	Selection	WHERE
ρ	Rename	AS

Relational Algebra Expression	SQL Statement
$\sigma_{\text{MaxSpeed} > 500}(\text{AIRPLANE})$	<pre>SELECT * FROM Airplane WHERE MaxSpeed > 500</pre>
$\pi_{P.Name}(\rho_P(\text{PILOT}))$	<pre>SELECT DISTINCT P.Name FROM Pilot AS P</pre>
$\pi_{\text{Date, Runway}}(\sigma_{\text{Date} \leq 9/1/2017 \text{ AND Runway} \neq \text{'W-3'}}(\text{FLIGHT}))$	<pre>SELECT DISTINCT Date, Runway FROM Flight WHERE Date <= 9/1/2017 AND Runway <> 'W-3'</pre>

Relational Algebra Expression	SQL Statement
$RANK \times SUI$	<pre>SELECT * FROM `Rank`, Suit; SELECT * FROM `Rank` CROSS JOIN Suit;</pre>
$FLIGHT \bowtie PILOT$	<pre>SELECT * FROM Flight NATURAL JOIN Pilot</pre>
$U \bowtie_{A < V.C} V$	<pre>SELECT DISTINCT * FROM U, V WHERE A < V.C; -- or, using "infix" join syntax: SELECT DISTINCT * FROM U INNER JOIN V ON A < V.C;</pre>

The use of literal dates in SQL requires care.

Example (*as-is, this SQL statement will not produce the expected result*):

$\sigma_{\text{Date} \leq 9/1/2017 \text{ AND Runway} \neq \text{'W-3'}}(\text{FLIGHT})$

```
SELECT *  
FROM Flight  
WHERE Date <= 9/1/2017  
      AND Runway != 'W-3'
```

- Option 1: Use string literals in a recognized format
 - ANSI: 'YYYY-MM-DD HH:MM:SS'
 - [ISO 8601](#) date format allows you to specify time zone (ie. '2017-10-05T17:47:17Z')
 - RDBMSs often have other default string representations of dates (typically, the default is configurable, so it's dangerous to rely on this!)
- Option 2: Use vendor-specific date functions
 - MySQL
 - [MAKEDATE\(<year>, <day of year>\)](#)
 - [STR_TO_DATE\(<string>, <format>\)](#) (format is C-style: %m, %d, %Y, etc.)

A date value (literal or a table/column reference) may be used along with [MySQL date functions](#) to represent expressions such as: "within the last week", "last year", "how many days since ...?", etc.

```
SELECT DATE_ADD(DateEnrolled, INTERVAL 4 YEAR) AS TargetGradDate  
FROM Student;
```

```
SELECT DateEnrolled, DATEDIFF(CURRENT_TIMESTAMP, DateEnrolled) / 365 AS YearsEnrolled  
FROM Student  
WHERE StudentID = '146461564';
```

Due to the use of vendor-specific functions, these two SQL statements are *MySQL-specific*, not ANSI SQL!

Revised examples:

```
SELECT *  
FROM Flight  
WHERE Date = '2019-10-06';
```

```
SELECT *  
FROM Flight  
WHERE Date <= STR_TO_DATE('9/1/2019', '%m/%d/%Y') AND Runway != 'W-3';
```

```
SELECT *  
FROM Flight  
WHERE Date >= DATE_SUB(CURRENT_DATE, INTERVAL 4 MONTH); -- within the last 4 months
```

We will encounter a few more scalar functions in SQL, and will introduce them as needed.

There is little standardization in this area. For the most part, each RDBMS vendor supports its own unique collection of scalar functions.

	Relational Algebra	SQL
Set Union	\cup	UNION
Set Difference	$-$	EXCEPT
Set Intersection	\cap	INTERSECT

- SQL defines two UNION variants:
 - UNION - Uses set rules (eliminates duplicates)
 - UNION ALL - Bag/multiset rules (preserves duplicate rows)
- Just as in relational algebra, SQL set operators (including UNION and UNION ALL) require *union compatibility*:
 - Same number of columns and matching data types.


```
(SELECT Name FROM Pilot)  
UNION  
(SELECT Make FROM Airplane)
```

```
(SELECT Name FROM Pilot)  
UNION ALL  
(SELECT Make FROM Airplane)
```

TaxID	FirstName	LastName
458-60-6366	Helen	Medina
534-42-0424	Ann	Mills

TaxID	FirstName	LastName
549-81-3606	Gary	Russell
458-60-6366	Helen	Medina

UNION

**UNION
ALL**

TaxID	FirstName	LastName
458-60-6366	Helen	Medina
534-42-0424	Ann	Mills
549-81-3606	Gary	Russell

TaxID	FirstName	LastName
458-60-6366	Helen	Medina
534-42-0424	Ann	Mills
549-81-3606	Gary	Russell
458-60-6366	Helen	Medina

INTERSECT returns only rows that are present in *both* result sets.

-- List names of pilots that are also airplane makers

```
SELECT Name FROM Pilot
```

```
INTERSECT
```

```
SELECT Make as Name FROM Airplane
```

`EXCEPT` returns rows that are present in the first query, but not present in the second query. Some vendors use the keyword `MINUS`

```
-- List IDs of pilots who do not have any flight records
SELECT PilotId FROM Pilot
EXCEPT
SELECT PilotId FROM Flight
```

- UNION, EXCEPT, and INTERSECT use *set semantics* by default, but the ANSI standard defines an ALL qualifier that preserves duplicates (bag or multiset semantics)
- Selection works the same on bags as it does on sets
- Projection with bag semantics *does not* eliminate duplicates
 - SELECT : bag semantics, no duplicate elimination!
 - SELECT DISTINCT : eliminates duplicate tuples
- Cartesian Product and joins operate on pairs of tuples. No difference in how they work when considering sets vs bags.

Adding the `DISTINCT` keyword at the start of the `SELECT` clause causes only unique tuples to be returned (in other words: when you include `DISTINCT`, the result is a *set* of tuples rather than a multiset / bag)

```
SELECT DISTINCT Name  
FROM Pilot
```

```
SELECT DISTINCT Make, Model  
FROM Airplane
```

- Join operators in relational algebra
 - Natural join
 - Theta join
 - Equijoin
 - Semijoin
 - Antijoin
- Selectively pair tuples from two relations.
- SQL supports the join operators listed above (*and a few others*)

There are two ways to express cartesian product in SQL.

Implicit (*often accidental!*):

```
SELECT *  
FROM Pilot, Flight
```

Explicit:

```
SELECT *  
FROM Pilot CROSS JOIN Flight
```


There are multiple ways to express a natural join in SQL.

Explicit projection and equijoin on common column(s):

```
SELECT DISTINCT Flight.*, Pilot.Name
FROM Flight, Pilot
WHERE Pilot.PilotID = Flight.PilotID
```

Implicitly, using NATURAL JOIN

```
SELECT DISTINCT *
FROM Flight NATURAL JOIN Pilot
```

Why would we prefer one over the other?

Consider two relations:

EMP(EmpID, FirstName, LastName,
DeptID, Building, RoomNum)

DEPT(DeptID, Name, Building)

The SQL standard defines `JOIN ... USING` as a shortcut for natural equijoins. This allows explicit control over the join columns (vs. `NATURAL JOIN`)

```
SELECT *  
FROM Emp INNER JOIN Dept USING (DeptID)
```

Consider the same two relations:

EMP(EmpID, FirstName, LastName,
 DeptID, Building, RoomNum)

DEPT(DeptID, Name, Building)

We can use the following form to express theta joins:

Relational algebra expression:

$$U \bowtie_{A < V.C \text{ AND } U.B \neq V.B} V$$

```
SELECT *  
FROM U, V  
WHERE A < V.C AND U.B <> V.B
```

ANSI SQL defines <> as the "not equal to" operator
Most RDBMSs support either <> or !=

As an alternative, we can use the following to express theta joins. Note the lack of a `WHERE` clause. Instead, we use the ANSI ***infix*** operator, `[INNER] JOIN ... ON,` to specify our join condition. The `INNER` keyword is optional.

```
SELECT *  
FROM U INNER JOIN V ON A < V.C AND U.B <> V.B
```

Relational algebra expression (same as previous slide):

$$U \bowtie_{A < V.C \text{ AND } U.B \neq V.B} V$$

Two ways to express theta join:

```
SELECT * FROM U, V WHERE A < V.C AND U.B <> V.B
```

```
SELECT * FROM U INNER JOIN V ON A < V.C AND U.B <> V.B
```

Same condition



- What's the difference?

- JOIN ON can help prevent accidental cartesian products
- Readability, especially with more than one join (separates relationship/join logic from selection logic)
- Other types of joins we'll see in SQL (OUTER)

**Recommendation: JOIN ... ON for conditions used to pair tuples,
WHERE clause for conditions that "filter" tuples**

A table may be joined to itself (*table aliases required*)

Find hard-disk sizes that occur in two or more PCs.

$$\pi_{hd}(\sigma_{PC1.model \neq PC2.model \wedge PC1.hd = PC2.hd}(\varrho_{PC1}(PC) \times \varrho_{PC2}(PC)))$$

```
SELECT DISTINCT pc1.hd
FROM pc AS pc1, pc AS pc2
WHERE pc1.model <> pc2.model AND pc1.hd = pc2.hd

-- JOIN ... ON syntax remains an option
```

Joins and unions both combine data from multiple tables, but they do so in different ways.

Informally, joins combine *columns* while unions operate on entire *rows*.

Important to keep in mind behavior with empty tables.

Set Operations vs NATURAL JOIN

TaxID	FirstName	LastName	StartDate
458-60-6366	Helen	Medina	2019-04-01
534-42-0424	Ann	Mills	2018-01-01

TaxID	FirstName	LastName
549-81-3606	Gary	Russell
458-60-6366	Helen	Medina

NATURAL JOIN

TaxID	FirstName	LastName	StartDate
458-60-6366	Helen	Medina	2019-04-01

UNION

Error: not union compatible

A single SELECT statement may include any number of JOINS:

```
SELECT Pilot.Name, Airplane.Make, Airplane.Model, Airplane.TailNum, Flight.Date
FROM Pilot
    INNER JOIN Flight ON Pilot.PilotID = Flight.PilotID
    INNER JOIN Airplane ON Airplane.TailNum = Flight.TailNum
```

For inner joins, the join order has no impact on results.

A SELECT statement may combine JOIN and WHERE (and a few more clauses that we will discuss shortly):

```
SELECT DISTINCT P.Name AS Pilot, A.Make, A.Model, A.TailNum, F.Date
FROM Pilot P
    INNER JOIN Flight F ON P.PilotID = F.PilotID
    INNER JOIN Airplane A ON A.TailNum = F.TailNum
WHERE F.Date <= '2019-12-31' AND A.Make <> 'Boeing' AND MaxSpeed > 500
```

Corresponding Relational Algebra expression?

Logical operators in SQL adhere to the following precedence rules:

AND takes precedence over OR, NOT takes precedence over both

Parentheses are advised. **DeMorgan's laws** are handy to keep in mind when dealing with complex expressions:

$$\text{NOT (A OR B)} = \text{NOT (A) AND NOT (B)}$$

$$\text{NOT (A AND B)} = \text{NOT (A) OR NOT (B)}$$

ANSI SQL supports a simple string matching operator:

```
<string> LIKE <pattern>
```

<pattern> may include wildcards: % (zero or more characters) or _ (one character)

```
SELECT *  
FROM Pilot  
WHERE Name LIKE 'Vic%'
```

Also:

```
<string> NOT LIKE <pattern>
```

```
SELECT *  
FROM Pilot  
WHERE Name NOT LIKE '%_r'
```

Warning: In MySQL, LIKE and NOT LIKE are **case INsensitive** by default!

Some RDBMSs allow more advanced pattern matching (for example: [MySQL REGEX](#) / [NOT REGEX](#)) However, this is not part of the ANSI standard.

NULL handling requires care in SQL. Specialized comparison operators:

IS NULL

IS NOT NULL

SELECT *

FROM Airplane

WHERE MaxSpeed IS NOT NULL

None of these will work as expected:

WHERE MaxSpeed = NULL

WHERE NOT (MaxSpeed = NULL)

WHERE MaxSpeed != NULL

```
SELECT *  
FROM Airplane  
WHERE MaxSpeed > 250 OR MaxSpeed <= 250
```

```
SELECT *  
FROM Airplane  
WHERE (MaxSpeed = 220) OR NOT (MaxSpeed = 220)
```

x	y	x AND y	x OR y	NOT x
TRUE	TRUE	TRUE	TRUE	FALSE
TRUE	UNKNOWN	UNKNOWN	TRUE	FALSE
TRUE	FALSE	FALSE	TRUE	FALSE
UNKNOWN	TRUE	UNKNOWN	TRUE	UNKNOWN
UNKNOWN	UNKNOWN	UNKNOWN	UNKNOWN	UNKNOWN
UNKNOWN	FALSE	FALSE	UNKNOWN	UNKNOWN
FALSE	TRUE	FALSE	TRUE	TRUE
FALSE	UNKNOWN	FALSE	UNKNOWN	TRUE
FALSE	FALSE	FALSE	FALSE	TRUE

To control the order in which tuples are returned, we may append the optional `ORDER BY` clause to a `SELECT` statement. `ORDER BY` may include one or more attributes.

```
SELECT *  
FROM Airplane  
ORDER BY Make, Model
```

By default, order is lowest to highest for each attribute. We can reverse the order (per attribute) by appending `DESC` to any attribute.

SQL supports a variety of expressions, including basic arithmetic and string/date scalar functions. Examples:

```
SELECT Make, Model, MaxSpeed * 1.60934 AS SpeedKPH  
FROM Airplane
```

```
SELECT Name, CHAR_LENGTH(Name) AS NameLength  
FROM Pilot  
ORDER BY CHAR_LENGTH(Name)
```

`CHAR_LENGTH()` counts characters — taking into account multi-byte character sets — `LENGTH()` counts *bytes*

More examples:

CONCAT() and DATE_SUB()
are *MySQL-specific* functions!

```
SELECT CONCAT(Make, ' ', Model, '/', TailNum) AS PlaneDescription
FROM Airplane
WHERE CONCAT(Make, ' ', Model) LIKE '%Air%' OR MaxSpeed IS NULL
```

```
SELECT Date, DATE_SUB(Date, INTERVAL 1 DAY) AS FlightPlanDue
FROM Flight
WHERE Date <= CURRENT_DATE
ORDER BY Date DESC
```

We have seen basic comparison and boolean connectors in the `WHERE` clause. SQL also supports the following optional syntax for convenience:

SQL Condition	Equivalent to...
<code>column BETWEEN x AND y</code>	<code>column >= x AND column <= y</code>
<code>column NOT BETWEEN x AND y</code>	<code>column < x OR column > y</code>
<code>column IN (a, b, c)</code>	<code>column = a OR column = b OR column = c</code>
<code>column NOT IN (a, b, c)</code>	<code>column <> a AND column <> b AND column <> c</code>

3	SELECT
4	DISTINCT
1	FROM ... JOIN ... ON
2	WHERE
5	UNION / INTERSECT / EXCEPT
6	ORDER BY

SQL	Relational Algebra Operator	
SELECT DISTINCT	π	Projection
FROM	\times	Cartesian product
INNER JOIN ... ON	\bowtie_{θ}	Theta join
WHERE	σ	Selection
AS	ρ	Rename
UNION	\cup	Set Union
EXCEPT	$-$	Set Difference
INTERSECT	\cap	Set Intersection

- The SQL language fully supports the five core relational algebra operators and derived operators (joins, intersect). We will introduce a few additional tools, including:
 - Additional join variations
 - Grouping / aggregation (GROUP BY)

The WHERE clause can be used with SELECT (as we have seen). The same WHERE clause syntax applies to UPDATE, and DELETE.

```
SELECT *  
FROM Flight  
WHERE Date > CURRENT_DATE AND Runway LIKE 'W%';
```

```
DELETE  
FROM Flight  
WHERE Date > CURRENT_DATE AND Runway LIKE 'W%';
```


The UPDATE statement can be used to change multiple columns in multiple rows of a *single* table.

Syntax:

```
UPDATE <table>  
SET <column1> = <value1> [, <column2> = <value2>]  
[WHERE <predicate>];
```

All flights that depart on a future date from a runway that begins with "W" now depart on the following day.

```
UPDATE Flight
SET Date = DATE_ADD(Date, INTERVAL 1 DAY)
WHERE Date > CURRENT_DATE AND Runway LIKE 'W%';
```

A recent software update has increased the maximum speed of all Boeing aircraft by 10%. Also, all Boeing model numbers are now suffixed with an "s"

```
UPDATE Airplane  
SET MaxSpeed = MaxSpeed * 1.1, Model = CONCAT(Model, 's')  
WHERE Make = 'Boeing';
```

All airplanes must undergo regular inspections. Add a new column to the Airplane table to record the last date of inspection. Set this value to January 29, 2024 for all airplanes.

```
ALTER TABLE Airplane ADD COLUMN LastInspected DATE;
```

```
UPDATE Airplane  
SET LastInspected = '2023-04-26';
```

No WHERE clause, **all** rows
will be updated!

In addition to the inspection date, we need to store the initials of the mechanic who performed each inspection, as well as any notes he/she may have recorded.

To [add a foreign key constraint](#):

```
ALTER TABLE <table>
    ADD [CONSTRAINT [name of constraint]]
    FOREIGN KEY (<column>, ...)
    REFERENCES <table> (<column>, ...);
```

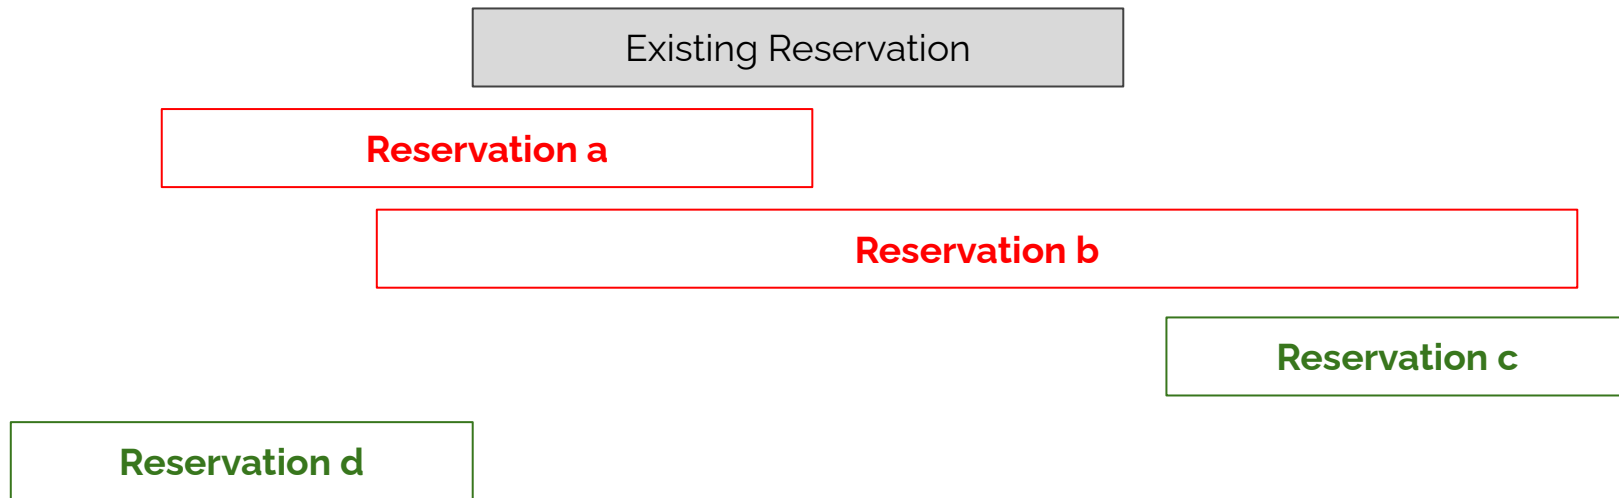
To drop a column:

```
ALTER TABLE <table> DROP [COLUMN] <column>;
```

Given a relation with schema:

RESERVATIONS(Code, Room, CheckIn, CheckOut)

How could we identify reservations that would overlap for a given date range?
(ie. where at least one night would be double-booked for the room)



To avoid overlap, a reservation must (for every existing reservation) start after the "other" reservation ends or end before the "other" reservation begins.