

# CSC 365

## Introduction to Database Systems

SQL	Relational Algebra Operator	
SELECT DISTINCT	$\pi$	Projection
FROM	$\times$	Cartesian product
INNER JOIN ... ON	$\bowtie_{\theta}$	Theta join
WHERE	$\sigma$	Selection
AS	$\rho$	Rename
UNION	$\cup$	Set Union
EXCEPT	$-$	Set Difference
INTERSECT	$\cap$	Set Intersection

Syntax:

```
SELECT <list of columns>  
FROM <table(s)>  
[ INNER|CROSS|LEFT|RIGHT|FULL JOIN <table> ON <join condition> ]  
[ WHERE <predicate> ]  
[ GROUP BY <columns> ]  
[ HAVING <group filter> ]  
[ ORDER BY <column(s)>
```

We have seen **scalar** functions in SQL that operate on a single row and return just one value. Previous examples:

```
SELECT CONCAT(Make, ' ', Model, '/', TailNum) AS PlaneDescription
FROM Airplane
WHERE CONCAT(Make, ' ', Model) LIKE '%Air%' OR MaxSpeed IS NULL
```

```
SELECT Date, DATE_SUB(Date, INTERVAL 1 DAY) AS FlightPlanDue
FROM Flight
WHERE Date <= CURRENT_DATE
ORDER BY Date DESC
```

In addition to scalar functions, SQL SELECT also supports **Aggregate Functions**, which summarize multiple rows.

- COUNT ( )
- MIN ( )
- MAX ( )
- SUM ( )
- AVG ( )

Simple examples:

```
SELECT COUNT(*) FROM Student      -- total number of students
```

With a WHERE clause:

```
SELECT COUNT(*) FROM Student      -- number of CSC majors
WHERE MajorCode = 'CSC'
```

Returns the lowest value from a single column.

```
SELECT MIN(DateEnrolled) FROM Student
```

May also be combined with WHERE:

```
SELECT MIN(DateEnrolled) FROM Student  
WHERE MajorCode IN ('CSC', 'SE', 'CPE')
```

MAX ( ) returns the highest value from a single column.

```
SELECT MAX(MaxSpeed)
FROM Airplane
```

It's valid to include multiple aggregates in a single SELECT:

```
SELECT MAX(MaxSpeed) , MIN(MaxSpeed) , COUNT (*)
FROM Airplane
```



SUM() and AVG() operate on numeric (integer or decimal) columns, returning the arithmetic sum or average.

```
SELECT SUM(MaxSpeed), COUNT(*), AVG(MaxSpeed)
FROM Airplane
```

Handling NULL values:

```
SELECT SUM(MaxSpeed) / COUNT(*) AS CalcAvg, AVG(MaxSpeed)
FROM Airplane
```

By default, SQL aggregate functions operate in "ALL" mode:

```
SELECT COUNT (Make) , COUNT (ALL Make) FROM Airplane
```

DISTINCT causes the aggregate function to consider only unique values:

```
SELECT COUNT (Make) , COUNT (DISTINCT Make) FROM Airplane
```

The functions AVG ( ) , COUNT ( ) , and SUM ( ) support DISTINCT mode

- Used alone, an aggregate function collapses all rows into *one* summarized row.
- What if we want to include other column values in our result set?


```
SELECT Make, AVG(MaxSpeed)  
FROM Airplane
```

The GROUP BY clause, along with aggregate functions, allows us to identify which column(s) to use as the basis for our summary:

```
SELECT Make, AVG(MaxSpeed)
FROM Airplane
GROUP BY Make
```

```
SELECT DAYNAME(DateEnrolled), COUNT(*)
FROM Student
GROUP BY DAYOFWEEK(DateEnrolled), DAYNAME(DateEnrolled)
ORDER BY DAYOFWEEK(DateEnrolled)
```

DAYNAME() and  
DAYOFWEEK() are  
*MySQL-specific* scalar date  
functions. [Full list here](#)



GROUP BY is permitted on both columns and the result of *scalar* functions:

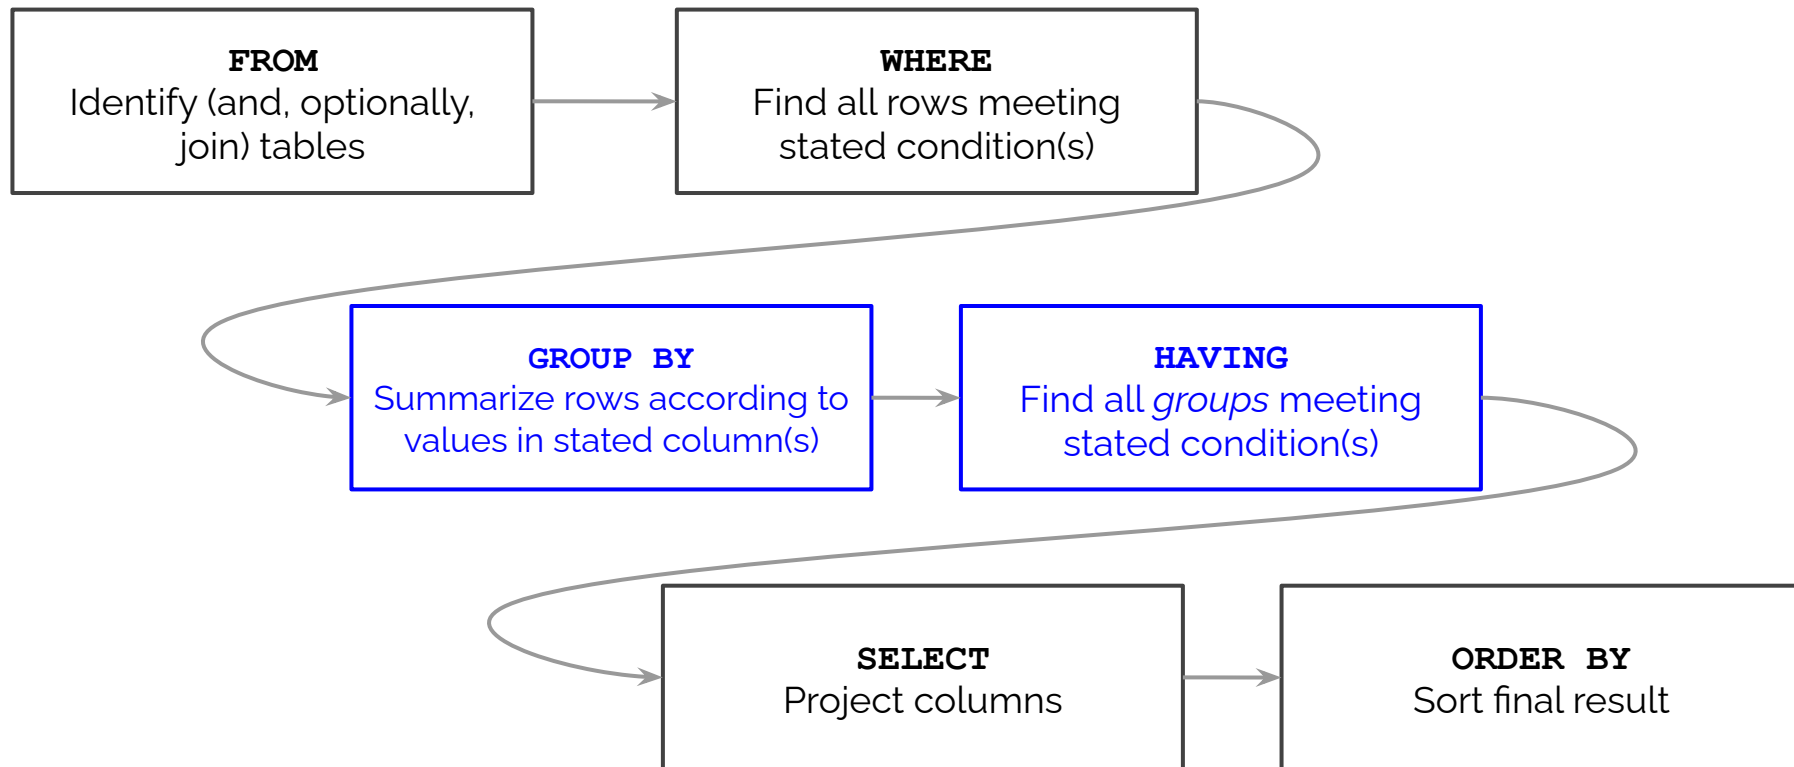
```
-- total guest count for each non-rustic room / month
SELECT RoomName, YEAR(CheckIn), MONTH(CheckIn),
       SUM(Adults + Kids) AS TotalGuests
FROM INN.rooms AS rm
     INNER JOIN INN.reservations AS res ON rm.RoomCode = res.Room
WHERE rm.decor NOT IN ('rustic')
GROUP BY RoomName, YEAR(CheckIn), MONTH(CheckIn)
```

Conditions on aggregate values require special treatment.

List all dates during October 2010 on which more than 5 people checked in to our inn. First try:

```
SELECT CheckIn, SUM(Adults + Kids) AS TotalGuests
FROM INN.reservations
WHERE CheckIn BETWEEN '2010-10-01' AND '2010-10-31'
      AND SUM(Adults + Kids) > 5
GROUP BY CheckIn
ORDER BY CheckIn
```

# SQL SELECT Logical Processing Order



Order	Clause
5	SELECT
6	DISTINCT
1	FROM ... JOIN ... ON
2	WHERE
3	GROUP BY
4	HAVING
7	UNION
8	ORDER BY



The HAVING clause, used along with GROUP BY, allows us to apply conditions that involve aggregate functions.

List all dates during October 2010 on which more than 5 people checked in to our inn.

```
SELECT CheckIn, SUM(Adults + Kids) AS TotalGuests
FROM INN.reservations
WHERE CheckIn BETWEEN '2010-10-01' AND '2010-10-31'
GROUP BY CheckIn
HAVING SUM(Adults + Kids) > 5
ORDER BY CheckIn
```

Find all airports that are the source of at least three Southwest flights. Report just the three-letter codes of the airports each code exactly once, in alphabetical order.

```
SELECT f.Source
FROM flights f
      INNER JOIN airlines a ON f.Airline = a.Id
WHERE a.Abbbr = 'Southwest'
GROUP BY f.Source
HAVING COUNT(*) >= 3
ORDER BY f.Source
```

1. Find the average price of a Cookie.
2. List all customers who have spent more than \$100 at the bakery. Include the customer's total spending and the number of days between the customer's most recent purchase and the end of October 2007.
3. What was the *average receipt amount* on Fridays versus Mondays?
4. List first and last name(s) of customers who have purchased *every* type of Danish

The *grouping operator*  $\gamma$  combines the effect of grouping and aggregation.

$$\gamma_L(R)$$

L is a list that may contain any of the following:

- One or more attributes of R, which will serve as *grouping attributes*
- Aggregation function(s) with an arrow and new name for the aggregate value

Examples of the *grouping operator*  $\gamma$  in relational algebra

$\gamma_{\text{Make, MAX(MaxSpeed)} \rightarrow \text{SpeediestInFleet}}$  (AIRPLANE)

```
-- highest maximum speed for each airplane maker
SELECT Make, MAX(MaxSpeed) AS SpeediestInFleet
FROM Airplane
GROUP BY Make
```

$\gamma_{\text{CheckInDate, SUM(Adults + Kids) \rightarrow \text{TotalGuests}}$  ( $\sigma_{\text{CheckIn} \geq \text{'2010-10-02'} \wedge \text{CheckIn} \leq \text{'2010-10-31'}}$  (RESERVATIONS))

```
-- number of guests who checked in on each date
SELECT CheckIn, SUM(Adults + Kids) AS TotalGuests
FROM reservations
WHERE CheckIn BETWEEN '2010-10-20' AND '2010-10-31'
GROUP BY CheckIn
```

Some references use a calligraphic G (**G**) rather than lowercase gamma ( $\gamma$ ):

Make **G** MAX(MaxSpeed)  $\rightarrow$  SpeediestInFleet (AIRPLANE)

CheckInDate **G** SUM(Adults + Kids)  $\rightarrow$  TotalGuests ( $\sigma$  CheckIn >= '10/1/2010' (RESERVATIONS))

Two more "extended" relational algebra operators:

- The duplicate-elimination operator  $\delta$  turns a bag into a set by eliminating all but one copy of each tuple (in SQL: `SELECT DISTINCT * . . .`)
- The sorting operator  $\tau$  turns a relation into a *list* of tuples, sorted according to one or more attributes (*Sorting should be applied cautiously. Some relational algebra operators do not make sense on lists.*)

How would we represent the following query in relational algebra, applying our extended relational algebra operators, gamma ( $\gamma$ ) and tau ( $\tau$ )?

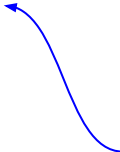
```
-- dates during October 2010 on which more than 5 people checked in
SELECT CheckIn, SUM(Adults + Kids) AS TotalGuests
FROM INN.reservations res
      JOIN INN.rooms r ON r.RoomCode = res.Room
WHERE CheckIn BETWEEN '2010-10-01' AND '2010-10-31'
GROUP BY CheckIn
HAVING TotalGuests > 5
ORDER BY CheckIn
```



# Can we do this in a single SQL statement?

BAKERY: List total sales amount for each food and flavor combination, along with the percent of sales for each flavor within its type of pastry (Twist, Meringue, Cake, Danish, etc.)

Flavor	Food	TotalSales	PctOfFoodSales
Almond	Twist	20.70	100
Chocolate	Meringue	25	60.8
Vanilla	Meringue	16.10	39.2
...	...	...	...



% should sum to 100 for each pastry type (Food column)

Window functions allow us to perform calculations across a result set. Somewhat like aggregation / GROUP BY, but window functions *do not* collapse rows. All rows are preserved.

Syntax (in SELECT clause):

<window function> OVER (PARTITION BY <column list> [ORDER BY <columns>])

Rows with the same value for  
<column list> fall into the same  
partition

Available window functions include  
familiar aggregate functions  
(SUM, MIN, MAX, COUNT, AVG)  
as well as a few new functions:  
RANK, ROW\_NUMBER, NTILE, LAG, LEAD

Function Name	Description
<a href="#">RANK()</a>	Rank of current row within its partition, <i>with gaps</i>
<a href="#">ROW_NUMBER()</a>	Number of current row within its partition
<a href="#">DENSE_RANK()</a>	Rank of current row within its partition, without gaps
<a href="#">FIRST_VALUE()</a>	Value of argument from first row of partition
<a href="#">LAST_VALUE()</a>	Value of argument from last row of partition
<a href="#">NTH_VALUE()</a>	Value of argument from N-th row of window frame
<a href="#">NTILE()</a>	Bucket number of current row within its partition.

[\(plus a few more...\)](#)

```
-- All employees, along with department average salary
SELECT DeptID, EmpID, AnnualSalary,
       AVG(AnnualSalary) OVER (PARTITION BY DeptID) as DeptAverage
FROM employee;
```

```
-- In which order were employees hired within their department?
SELECT DeptID, EmpID, HireDate,
       RANK() OVER (PARTITION BY DeptID ORDER BY HireDate) as HireOrder
FROM employee;
```

```
-- RANK() vs DENSE_RANK()
```

Shared, named windows can be defined using the `WINDOW` clause, which appears between the `HAVING` and `ORDER BY` clauses.


```
WINDOW <window name> AS (window spec)
      [, <window name> AS (window spec)] ...
```

```
SELECT DeptID, EmpID, HireDate,  
       RANK() OVER w AS HireOrder,  
       DENSE_RANK() OVER w AS DenseHireOrder,  
       FIRST_VALUE(HireDate) OVER w AS EarliestDeptHire  
FROM employee  
WINDOW w AS (PARTITION BY DeptID ORDER BY HireDate)
```

```
-- How much less is each employee paid than the highest compensated
-- in his/her department
SELECT DeptID, EmpID, HireDate, AnnualSalary,
       RANK() OVER (w ORDER BY HireDate) AS HireOrder,
       AVG(AnnualSalary) OVER w AS DeptAvgSalary,
       MAX(AnnualSalary) OVER w AS DeptMaxSalary,
       (MAX(AnnualSalary) OVER w) - AnnualSalary AS SalaryDiff
FROM employee
WINDOW w AS (PARTITION BY DeptID)
```

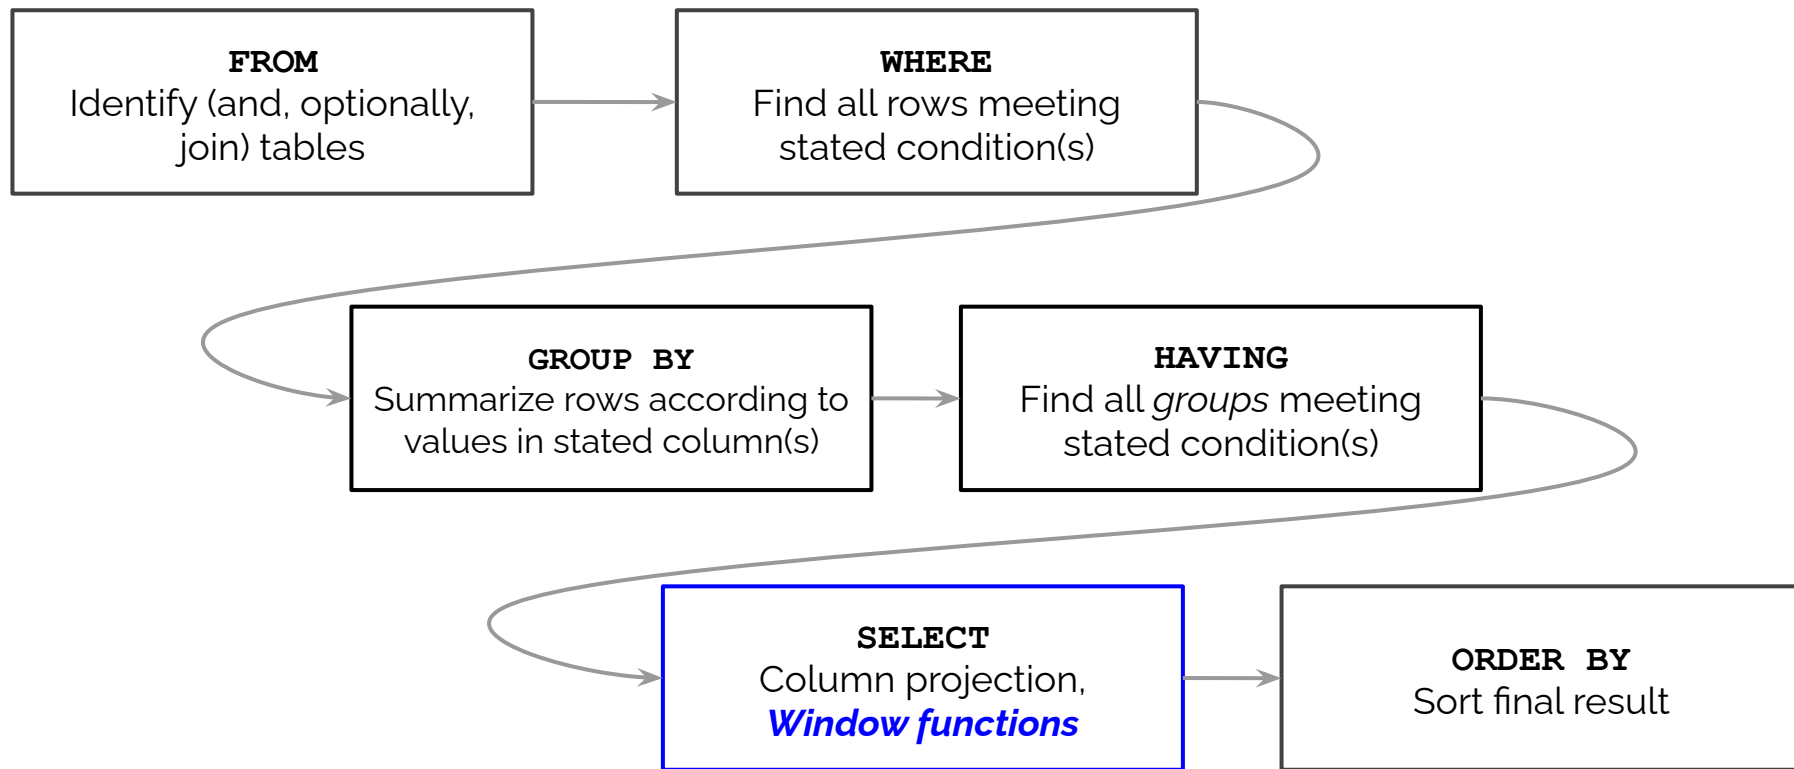
BAKERY: List total sales amount for each food and flavor combination, along with the percent of sales for each flavor within its type of pastry (Twist, Meringue, Cake, Danish, etc.)

Flavor	Food	TotalSales	PctOfFoodSales
Almond	Twist	20.70	100
Chocolate	Meringue	25	60.8
Vanilla	Meringue	16.10	39.2
...	...	...	...



% should  
sum to 100  
for each  
pastry type  
(a.k.a. food)





- Window functions may be used only in the `SELECT` and `ORDER BY` clauses (recall logical query processing order)
- Window functions execute *after* `GROUP BY` and aggregate functions.
  - Valid to use an aggregate function call in the arguments of a window function
  - Not possible to use window functions within regular aggregate functions
- Not supported by all RDBMSs / versions
  - Available only since MySQL v8.0
  - No support in SQLite

**Online Transaction Processing (OLTP):** emphasis on high throughput, inserts/updates, speed, concurrency, and recoverability.

**Online Analytical Processing (OLAP):** complex queries and analysis (often involving aggregation), sometimes referred to as decision-support queries.

OLTP	OLAP
<pre>INSERT INTO employee (EmpID, DeptID, Salary, FirstName, LastName, HireDate) VALUES (...);  SELECT * FROM employee WHERE HireDate &gt; '2018-01-01' AND DeptID IN ('ENG', 'HR', 'ACCTG');  UPDATE employee SET Salary = Salary * 1.1 WHERE HireDate BETWEEN '2017-01-01' AND '2017-03-31' AND Salary &lt; 50000;</pre>	<pre>SELECT AVG(Salary),       MIN(Salary),       MAX(Salary) FROM employee;  SELECT DeptID, COUNT(EmpID) FROM employee WHERE DeptID &lt;&gt; 'SALES' GROUP BY DeptID HAVING AVG(Salary) &gt; 50000;</pre>

Low *selectivity* (few rows accessed at a time)  
high *projectivity* (number of columns accessed)

Low *projectivity*, high *selectivity*

Online Analytical Processing (OLAP) often benefits from column-oriented storage: all values in a particular column are stored contiguously.

EmpID	Department	Salary
001	HR	60000
002	ENG	80000
003	ENG	85000
004	SALES	120000

Online Transaction Processing (OLTP) typically uses row storage: data for each tuple/record stored together.

001	HR	60000
001	ENG	80000
001	ENG	85000
001	SALES	120000