

# CSC 365

## Introduction to Database Systems

Possible results from a SQL `SELECT`:

- **Empty set**
- **Single scalar value** - Can be included in the `SELECT` list. Also may be used for comparison in the `WHERE` clause
- **Relation / Table**
  - 1-Dimensional for `WHERE IN / ANY / ALL / EXISTS`
  - 2-Dimensional may appear in the `FROM` clause as a stand-in for a named table


## Example: Departments Established Before the Physics Dept.

---

```
-- find the date on which the Physics department was established
SELECT DateEstablished
FROM Department
WHERE Code = 'PHYS';
```

```
-- find departments established before the Physics department
SELECT *
FROM Department
WHERE DateEstablished < value returned by query above;
```

```
-- Departments established before the Physics department
SELECT *
FROM Department
WHERE DateEstablished < (
    SELECT DateEstablished
    FROM Department
    WHERE Code = 'PHYS'
)
```



This subquery returns one row  
/ one column, due to WHERE  
on primary key column.

-- can we do this without a subquery? (as in Lab 4)

Often, a query can be expressed using either a subquery or a JOIN. Another way to write the previous query:

```
SELECT D1.*  
FROM Department D1  
    INNER JOIN Department D2 ON D1.DateEstablished < D2.DateEstablished  
WHERE D2.Code = 'PHYS'
```

- Why choose one over the other?
  - Clarity, especially in complex queries
  - Accidental row duplication or elimination with JOINS
  - Performance (based on profiling)

Subqueries may be used with the following comparison operators in the `WHERE` clause:

- `IN`
- `ALL`
- `ANY`
- `EXISTS`

Each of these may be negated with `NOT` (`NOT IN`, `NOT EXISTS`, etc.)

- $s \text{ IN } R$  is true if and only if  $s$  is equal to one of the values in  $R$ .
  - $s$  may be a single value or single column name, in which case  $R$  must be a single-column relation
  - If  $s$  has more than one element, the number of elements must match the number of columns in  $R$
- $s \text{ NOT IN } R$  is true if any only if  $s$  is equal to *no value* in  $R$

```
-- Students without a minor who are in a department other than CSSE
SELECT *
FROM Student
WHERE MinorCode IS NULL
AND MajorCode NOT IN (
    SELECT Code
    FROM Discipline
    WHERE Dept = 'CSSE'
)
```

How does this query change if we use JOIN rather than a subquery?



$s$  <comparator> ALL  $R$

$s$  <comparator> ANY  $R$

(comparator may be any of the following: >, <, >=, <=, =, <>)

$s$  > ALL  $R$  is true if and only if  $s$  is greater than every value in *unary* relation  $R$ .

$s$  <> ALL  $R$  is the same as  $s$  NOT IN  $R$

$s$  > ANY  $R$  is true if and only if  $s$  is greater than at least one value in *unary*

relation  $R$ .  $s$  = ANY  $R$  is the same as  $s$  IN  $R$

ALL / ANY may be negated like other boolean expressions:

NOT  $s \geq \text{ALL } R$  is true if and only if  $s$  is not the maximum value in  $R$

NOT  $s > \text{ANY } R$  is true if any only if  $s$  is the minimum value in  $R$

```
-- Find departments established before every department in Engineering
SELECT *
FROM Department
WHERE DateEstablished < ALL (
    SELECT DateEstablished
    FROM Department WHERE College = 'CENG'
)
```

```
-- Find the earliest-established department
SELECT *
FROM Department
WHERE NOT DateEstablished > ANY (
    SELECT DateEstablished
    FROM Department
)
```

Demo Area:

- SUPPLIERS-DEMO-2
- OUTER-JOIN-*n* exercises

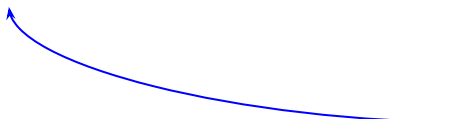
The subquery examples seen thus far have been uncorrelated: subqueries that can be executed once for the entire outer `SELECT`.

It is also possible to refer to values in the outer query from within a subquery, causing the subquery to be re-evaluated multiple times during execution of the outer query.

This second form of subquery is called a **correlated subquery**

```
-- List of all students, along with the number of students  
-- who share their major
```

```
SELECT FirstName, LastName, MajorCode,  
       (SELECT COUNT(*) FROM Student AS s2  
        WHERE s2.MajorCode = s.MajorCode) as TotalInMajor  
FROM Student AS s
```



Inner query references  
column value from outer  
query.

```
-- List sales detail, along with running total for each customer
SELECT C.Name, R.Item, R.PDate, P.Price AS ItemPrice,
       (SELECT SUM(P2.Price)
        FROM Purchases R2
        INNER JOIN Items P2 ON P2.SKU = R2.Item
        WHERE R2.Customer = R.Customer
        AND R2.PDate <= R.PDate) as CustomerRevenueToDate
FROM Purchases R
     INNER JOIN Customers C ON R.Customer = C.ID
     INNER JOIN Items P ON P.SKU = R.Item
WHERE C.Name <> 'Ben Vilasec'
ORDER BY Name, PDate
```

What does this query look in relational algebra?



The previous examples demonstrate **correlated subqueries** in the `SELECT` clause

Also permitted in the `WHERE` clause:

- Direct comparisons with scalar values
- `EXISTS` operator

- EXISTS  $R$  is true if and only if  $R$  is not empty
- NOT EXISTS  $R$  is true if and only if  $R$  is empty

```
-- Find students who are the only students in their major
SELECT *
FROM Student o
WHERE MajorCode IS NOT NULL
AND NOT EXISTS (
    SELECT StudentID
    FROM Student i
    WHERE o.MajorCode = i.MajorCode AND o.StudentID <> i.StudentID
)
```

```
-- Students who share a major with at least one other student?
```

In the WHERE clause, subqueries can be transformed into equivalent correlated subqueries using EXISTS & NOT EXISTS.

```
SELECT *  
FROM Student  
WHERE MajorCode IN (  
    SELECT Code  
    FROM Department  
    WHERE College = 'OCOB'  
)
```

```
SELECT *  
FROM Student  
WHERE EXISTS (  
    SELECT Code  
    FROM Department  
    WHERE College = 'OCOB'  
    AND Code = Student.MajorCode  
)
```

List students who were enrolled on/after all departments in CENG were established.

```
SELECT *  
FROM Student  
WHERE DateEnrolled >= ALL (  
    SELECT DateEstablished  
    FROM Department  
    WHERE College = 'CENG'  
)
```

```
SELECT *  
FROM Student  
WHERE NOT EXISTS (  
    SELECT DateEstablished  
    FROM Department  
    WHERE College = 'CENG'  
    AND Student.DateEnrolled < DateEstablished  
)
```



Relations R(A, B) and S(C)

```
SELECT C
FROM S
WHERE C IN (
    SELECT SUM(B)
    FROM R
    GROUP BY A
)
```

```
SELECT C
FROM S
WHERE EXISTS (
    SELECT SUM(B) FROM R
    GROUP BY A
    HAVING SUM(B) = S.C
)
```

*[From Homework 3]*

#9: Find all bands in which 'Irmin Schmidt' DID NOT play. Output the names of the bands.

#10: Find all 'Pink Floyd' band members who did NOT participate in the recording of the album 'Meddle'. (note: a musician did not participate in a recording of an album if he did not play in the band that year).

We've seen how IN, ANY, and ALL work with one-column relations. It is also possible to perform tuple-based comparisons using subqueries, as long as the *degree* matches on both sides of the comparison.

```
-- Students who share the same major & minor as another student
SELECT *
FROM Student SO
WHERE (MajorCode, MinorCode) IN (
    SELECT MajorCode, MinorCode FROM Student SI
    WHERE SI.StudentID <> SO.StudentID
)
-- Anybody missing?
```



Since a `SELECT` statement returns a relational table, we can use a nested `SELECT` statement in the `FROM` clause of a query.

```
SELECT <column list>  
FROM  (SELECT query) [AS] <alias>  
[WHERE <condition> ]  
[GROUP BY <attribute list>  
[HAVING <group condition>]
```

Two requirements to consider when using a nested `SELECT` statement in the `FROM` clause of a query:

1. Nested `SELECT` must be enclosed in parentheses and *must have an alias*
2. All computed columns (aggregates, scalar functions, `CASE`, etc.) *must have aliases*

(BAKERY dataset)

Show customer names, dates and total purchase amounts for days where the customer's total purchase amount for that day is more than two times the customer's average purchase amount on days when he/she made a purchase.

Total purchases by the customer (all time) divided by the number of days on which the customer made a purchase.

```
SELECT ?  
FROM BAKERY.customers C  
    INNER JOIN BAKERY.receipts R ON Customer = CId  
    INNER JOIN BAKERY.items ON RNumber = Receipt  
    INNER JOIN BAKERY.goods ON Item = GId  
GROUP BY ?
```

```
SELECT C.CId, FirstName, LastName, SaleDate,  
       ROUND(SUM(price), 2) AS DailyTotalPurch,  
       ROUND(P.AvgDailyPurchase, 2) AS AvgDailyPurchase  
FROM BAKERY.customers C  
     INNER JOIN BAKERY.receipts R ON Customer = CId  
     INNER JOIN BAKERY.items ON RNumber = Receipt  
     INNER JOIN BAKERY.goods ON Item = GId  
     INNER JOIN (  
         SELECT CId, SUM(price) / COUNT(DISTINCT SaleDate) AS AvgDailyPurchase  
         FROM BAKERY.customers, BAKERY.receipts, BAKERY.items, BAKERY.goods  
         WHERE CId = Customer AND RNumber = Receipt AND Item = GId  
         GROUP BY CId  
     ) P ON C.CId = P.CId  
GROUP BY CId, FirstName, LastName, SaleDate, P.AvgDailyPurchase  
HAVING SUM(price) > P.AvgDailyPurchase * 2.0;
```

Subqueries can be deeply nested (up to practical limits imposed by RDBMS implementations.) Inner queries can reference outer columns, but cannot reference siblings.


```
-- Non-CENG departments with students who enrolled
-- prior to the date the department was established
SELECT *
FROM Department d1
WHERE EXISTS (
    SELECT StudentID FROM Student s
    WHERE s.MajorCode = d1.Code AND DateEnrolled < (
        SELECT DateEstablished FROM Department d2
        WHERE d2.Code = s.MajorCode
    ) AND s.MajorCode NOT IN (
        SELECT Code FROM Department d3
        WHERE d3.College = 'CENG'
    )
)
-- How would we list the corresponding students?
```

- Considerable flexibility, deep nesting allowed
- Can appear throughout `SELECT` statement (with some exceptions)
- Subquery variations:
  - Uncorrelated
  - Derived table in `FROM` clause
  - Correlated
    - No correlation permitted for subqueries that appear `FROM` clause
- *Performance implications (profiling always recommended)*

ANSI SQL does not define a standard way to limit the number of rows returned by a `SELECT` query.

Each RDBMS has its own syntax. In MySQL, we use `LIMIT`:

```
SELECT *  
FROM Student  
ORDER BY LastName  
LIMIT 2
```



Sorting (`ORDER BY`) is performed **before** `LIMIT`, so we see the first two students, based on A-Z ordering of last name.



Single-argument `LIMIT` returns up to the specified number of rows:

```
LIMIT 5      -- top 5 rows in result set
```

Two-argument `LIMIT` returns up to the requested number of rows, after a given offset (zero indexed):

```
LIMIT 5,10   -- rows 6-15
```

RDBMS	Syntax
MySQL & PostgreSQL	<code>SELECT * FROM table LIMIT 10</code>
Microsoft SQL Server & Access	<code>SELECT TOP 10 * FROM table</code>
Oracle	<code>SELECT * FROM     (SELECT * FROM table) WHERE rownum &lt;= 10</code>
IBM DB2	<code>SELECT * FROM table FETCH FIRST 10 ROWS ONLY</code>
Informix	<code>SELECT FIRST 10 * FROM table</code>

```
SELECT Code, DeptName, DateEstablished,
```

```
    DATEDIFF((SELECT DateEstablished
```

```
        FROM Department
```

```
        ORDER BY DateEstablished DESC
```

```
        LIMIT 1),
```

```
    DateEstablished) / 365 AS YearsOlderThanNewestDept
```

```
FROM Department
```

Can we write this query  
without LIMIT?

```
-- List customer(s) with the fewest number of purchases
SELECT C.CId, C.FirstName, COUNT(*) AS PurchaseCount
FROM BAKERY.receipts R
     INNER JOIN BAKERY.customers C ON R.Customer = C.CId
GROUP BY C.CId
ORDER BY PurchaseCount
LIMIT 1
```