

CSC 365

Introduction to Database Systems

WITH allows you to define temporary result sets (tables) that exist for a single query. These auxiliary queries are referred to as a Common Table Expressions (CTEs).

```
WITH <CTE1 name> AS (  
    SELECT ...  
) [, <CTE2 name> AS (  
    SELECT ...  
) ]  
SELECT * FROM <table>, <CTE1 name>, <CTE2 name> ...
```

```
WITH monthly_sales AS (  
    SELECT monthname(TransDate) AS Month, SUM(Amount) AS Sales  
    FROM accounting_entries  
    WHERE AccountType = 'Sales'  
    GROUP BY monthname(TransDate)  
),  
monthly_cogs AS (                                -- COGS: Cost of Goods Sold  
    SELECT monthname(TransDate) AS Month, SUM(Amount) AS COGS  
    FROM accounting_entries  
    WHERE AccountType = 'COGS'  
    GROUP BY monthname(TransDate)  
)  
SELECT T.AccountType, T.Account,  
    TransDate, monthname(TransDate) as Month,  
    Amount,  
    C.COGS as MonthCOGSTotal,  
    S.Sales as MonthSalesTotal,  
    ROUND((T.Amount / S.Sales) * 100, 2) AS PctOfMonthlySales  
FROM accounting_entries T  
    INNER JOIN monthly_sales S ON monthname(TransDate) = S.Month  
    INNER JOIN monthly_cogs C ON monthname(TransDate) = C.Month  
ORDER BY TransDate, AccountType
```

Find the most popular flavor(s) of Eclair sold at the BAKERY, based on number of purchases.

Among the possible approaches:

1. Nesting
2. CTE (eliminate some redundancy)

- Can simplify certain queries, avoid redundancy
- Just like Window Functions, supported only in the latest version of MySQL
 - Good support in other popular RDBMSs
- Optimization "fence" in *some* cases
 - CTE query might be materialized as table, without any algebraic laws applied across the entire query
 - Consult DBMS documentation and query plan

WITH RECURSIVE allows a WITH query to refer to its *own* output. To sum the numbers from 1 to 100:

```
WITH RECURSIVE t(n) AS (  
    SELECT 1  
    UNION ALL  
    SELECT n+1 FROM t WHERE n < 100  
)  
SELECT sum(n) FROM t;
```

See Also:

- [Solving the Traveling Salesman Problem with Postgres Recursive CTEs](#)
- [SQL 3D Engine](#)

```
-- Enumerate all flights from airport AMW with <= 3 transfers
WITH RECURSIVE routes(the_source, the_dest, route, hops) as (
    SELECT Source as the_source, Destination,
    CAST(CONCAT(Source, '->', Destination) AS CHAR(150)) as route, 0 as hops
    FROM AIRLINES.flights
    WHERE Source = 'AMW'
    UNION
    SELECT routes.the_source, f.Destination,
    CONCAT(route, '->', f.Destination),
    routes.hops + 1 as hops
    FROM AIRLINES.flights f, routes
    WHERE f.Source = routes.the_dest and f.Destination <> 'AMW' AND hops < 3
)
SELECT * FROM routes
```

- Common Table Expressions ()
 - WITH
 - WITH RECURSIVE

Introduced with the SQL 2005 standard, added to MySQL in version 8.0, released several years ago

WITH allows you to define temporary result sets (tables) that exist for a single query. These auxiliary queries are referred to as a Common Table Expressions (CTEs).

```
WITH <CTE1 name> AS (  
    SELECT ...  
) , <CTE2 name> AS (  
    SELECT ...  
)  
SELECT * FROM <table> , <CTE1 name> , <CTE2 name> ...
```

```
with small_flightnos as (  
    select FlightNo, Source, Destination  
    from cte_flights  
    where FlightNo between 25 and 100  
),  
possible_dests AS (  
    select Name AS Airline, FlightNo, Source, Destination  
    from cte_flights  
        join cte_airlines on Airline = Id  
    where Destination in ('CVO', 'AED', 'ATO')  
)  
select b.*  
from small_flightnos a  
    join possible_dests b on (a.FlightNo = b.FlightNo)
```

[SQL Fiddle](#)

WITH RECURSIVE allows a WITH query to refer to its *own* output. Example to sum the numbers from 1 to 100:

```
WITH RECURSIVE t(n) AS (  
    SELECT 1  
    UNION ALL  
    SELECT n+1 FROM t WHERE n < 100  
)  
SELECT sum(n) FROM t;
```

See also: [Solving the Traveling Salesman Problem with Postgres Recursive CTEs](#)

WITH RECURSIVE allows a WITH query to refer to its *own* output. Example (from PostgreSQL documentation) to sum the numbers from 1 to 100:

```
WITH RECURSIVE t(n) AS (  
    VALUES (1)  
    UNION ALL  
    SELECT n+1 FROM t WHERE n < 100  
)  
SELECT sum(n) FROM t;
```

See also: [Solving the Traveling Salesman Problem with Postgres Recursive CTEs](#)

```
WITH RECURSIVE routes(the_source, the_dest, route, hops) as (  
    SELECT Source as the_source, Destination,  
    CAST(CONCAT(Source, '->', Destination) AS CHAR(150)) as route, 0 as hops  
    FROM cte_flights  
    WHERE Source = 'AMW'  
    UNION  
    SELECT routes.the_source, f.Destination,  
    CONCAT(route, '->', f.Destination),  
    routes.hops + 1 as hops  
    FROM cte_flights f, routes  
    WHERE f.Source = routes.the_dest and f.Destination <> 'AMW' AND hops < 3  
)  
SELECT * FROM routes
```

[SQL Fiddle](#)

[Virtual] View - Named SQL query whose results are not stored. Query will be run by the database as needed.

Materialized View - Query result stored as a physical table

- Virtual views are relations that are not physically stored like tables.
- Defined as a SQL `SELECT` statement
- Views can be queried as if they were regular tables.
- In certain cases, we can perform DML actions using virtual views

Syntax to define a new view:

```
CREATE VIEW <view name> AS <view definition>;
```

Remove a view with:

```
DROP VIEW <view name>;
```

Dropping a view *does not* affect any underlying data.

CREATE VIEW RoomSummary AS

```
SELECT Res.Code, Rm.RoomCode, LastName AS CustLastName,  
       CheckIn, CheckOut,  
       DATEDIFF(Checkout, CheckIn) AS Days,  
       DATEDIFF(Checkout, CheckIn) * Rate AS TotalCharge,  
       Adults + Kids AS GuestCount  
FROM INN.reservations Res  
      INNER JOIN INN.rooms Rm ON Res.Room = Rm.RoomCode  
ORDER BY DATEDIFF(Checkout, CheckIn) * Rate DESC
```

```
-- Rooms with total charge > $2,500
```

```
SELECT RoomName, BasePrice, RS.*
```

```
FROM INN.rooms R
```

```
    INNER JOIN RoomSummary RS ON RS.RoomCode = R.RoomCode
```

```
WHERE TotalCharge > 2500
```

```
ORDER BY TotalCharge DESC
```

A view may be updatable if defined using a "simple" SQL query. To allow this, a view must be defined using a limited set of SQL features. A few restrictions:

- `SELECT DISTINCT` not permitted
- `FROM` includes just one table
- `SELECT` list must include *all* required (`NOT NULL`) / key attributes

In other words, the view body must consist of an extremely simple query to permit DML statements "through" the view.

If a view is used frequently or is very expensive to compute (typical for OLAP queries), it is often useful to *materialize* the view (not supported in MySQL) This causes the query results to be stored.

Downsides:

- Storage usage
- Materialized view must be *maintained* by the RDBMS: Each time a change is made to an underlying base table, the materialized data must be updated.

- Incremental
 - Changes are propagated to view *as they occur* in base tables
 - Possible for simple views, difficult as views become more complex
 - Downside: not always wise for frequently-updated base data (performance considerations)
- Periodic Refresh
 - Entire materialized view is reconstructed periodically (ie. each night)
 - Downside: data in materialized view is out of date

Basic data access rules can be applied using SQL's "Data Control Language" statements:

- GRANT - Allow users/groups to perform certain actions on objects in the database
- REVOKE - Take away granted privileges

Basic syntax:

SELECT, INSERT, DELETE,
UPDATE, etc.



```
GRANT <privilege(s)> ON <object(s)> TO <user(s)>;
```

Examples:

Table or view name

```
GRANT SELECT, INSERT, UPDATE, DELETE ON reservations TO dclerk;  
GRANT SELECT ON reservations TO taccountant;
```

(see [documentation](#) for additional options in MySQL)

One common use of views, along with the `GRANT` command is to apply basic access control rules within a database. Examples:

- Managers should see detail about employees in their own department only
- An HR analyst may change any salary except their own or supervisors'
- A payroll accountant should see only payroll-related accounts/transactions

Some databases extend this concept to support finer-grained security rules (eg. [Oracle label security](#), [Postgres row level security](#))

- A virtual views is nothing more than a named, stored SQL query
 - Remove from database with: `DROP VIEW <view name>;`
 - Dropping a view does not affect any data in physical tables
 - Can be used, along with data control statements (`GRANT` / `REVOKE`) to enforce data access rules in a multi-user database
 - To see tables vs. views in MySQL: `show full tables;`
- Materialized views (where supported) can offer significant performance benefits, offset by the problem of *maintenance*
 - A materialized view caches results of complex (slow) queries
 - Changes to base tables must be propagated to materialized views
 - Storage costs must be considered

BAKERY: List total sales amount for each food and flavor combination, along with the percent of sales for each flavor within its type of pastry (Twist, Meringue, Cake, Danish, etc.)

Flavor	Food	TotalSales	PctOfFoodSales
Almond	Twist	20.70	100
Chocolate	Meringue	25	60.8
Vanilla	Meringue	16.10	39.2
...

% should
sum to 100
for each
pastry type
(a.k.a. food)