


Iterative Deepening

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
 - Run a DFS with depth limit 1. If no solution...
 - Run a DFS with depth limit 2. If no solution...
 - Run a DFS with depth limit 3.
- Isn't that wastefully redundant?
 - Generally, most work happens in the lowest level searched, so not so bad!




General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important ideas:
 - Fringe
 - Expansion
 - Exploration strategy
- Main question: which fringe nodes to explore first?

Cost-Sensitive Search




BFS finds the shortest path in terms of number of actions.
It does not find the least-cost path. We will now cover
a similar algorithm which does find the least-cost path.

41


Uniform Cost Search

Strategy: expand a
cheapest node on the
fringe first:

Fringe is a priority queue
(priority: cumulative cost)




Note: This is how
Dijkstra's algorithm
works but it finds all
shortest paths from
Start



42

Uniform Cost Search (UCS) Properties



- What nodes does UCS expand?
 - Processes all nodes with cost less than cheapest solution!
 - If that solution costs C^* and arcs cost at least ϵ , then the "effective depth" is roughly C^*/ϵ
 - Takes time $O(b^{C^*/\epsilon})$ (exponential in effective depth)
- How much space does the fringe take?
 - Has roughly the last tier, so $O(b^{C^*/\epsilon})$
- Is it complete?
 - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
 - Yes! (Proof next lecture via A*)



43

Uniform Cost Issues

- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
 - Explores options in every "direction"
 - No information about goal location
- We'll fix that soon!



44

UCS vs Dijkstra's algorithm: Differences

The algorithms are very similar but differ in the details
Think through the potential differences:

- What is the goal of each algorithm?
- Initializing the Queue?
- Size of Queue?
- Why is UCS more desirable in our context?

The One Queue

- All these search algorithms are the same except for fringe strategies
 - Conceptually, all fringes are priority queues (i.e., collections of nodes with attached priorities)
 - Practically, for DFS and BFS, you can avoid the $\log(n)$ overhead from an actual priority queue, by using stacks and queues
 - Can even code one implementation that uses a variation on a single interface



Informed Search


Can we do better than BFS, DFS, or UCS? **What if we have some information about where the goal is?** Up to now we have been searching in all directions and have no idea how far it is to the goal.

Heuristics are a procedure or process to achieve something that is not guaranteed to achieve the best outcome.


- Use a heuristic function to guide us toward the Goal
 - Greedy Best-first search
 - A* Search

Search Heuristics

- A heuristic (in Search Problems) is:
 - A function that *estimates* how close a state is to a goal
 - Designed for a particular search problem
 - Pathing or path finding?
 - Examples: Manhattan distance, Euclidean distance for pathing




Manhattan Distances



49



Greedy Best First Search



50

Greedy Search

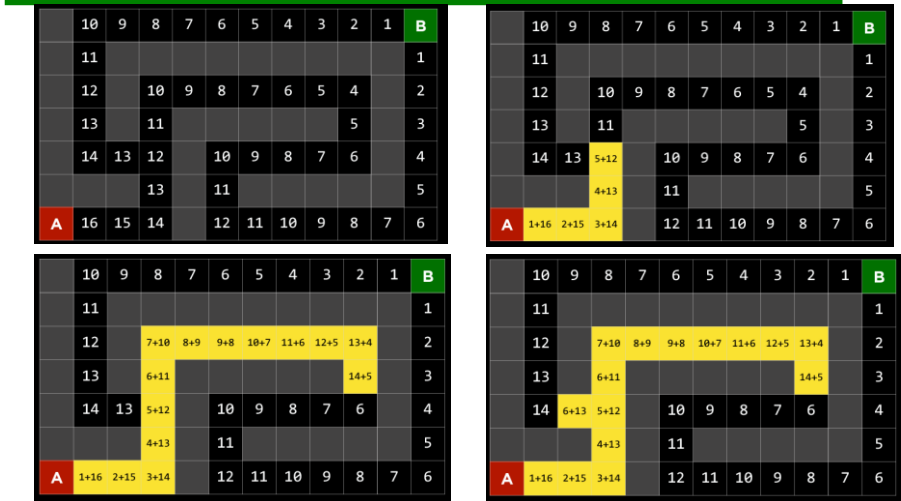
- Strategy: expand a node that you think is closest to a goal state
 - Heuristic: estimate of distance to nearest goal for each state
- A common case:
 - Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS



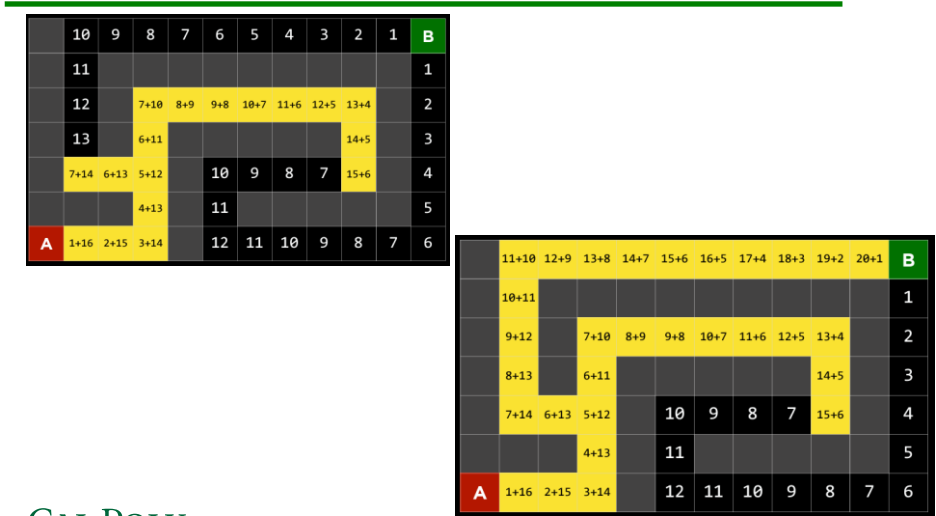
A* Search: Combine how far *traveled* with heuristic of *how far to get to goal*

- Search algorithm that expands node with lowest value of $f(n) = g(n) + h(n)$
 - $g(n)$ = cost to reach node (n)
 - $h(n)$ = estimated cost node to goal
(heuristic \approx what we used in greedy best-first search)

A* Search: Manhattan Distance Heuristic




A* Search: Manhattan Distance Heuristic




Properties of A*

Uniform-Cost




A*



55

When should A* terminate?

- Should we stop when we enqueue a goal?




- No: only stop when we dequeue a goal**

56

Is A* Optimal?

- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!





Admissible Heuristics

- A heuristic h is admissible (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal

- Examples:



0.0

- Coming up with admissible heuristics is most of what's involved in using A* in practice.


Optimality of A* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim:


- A will exit the fringe before B



Optimality of A* Tree Search

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A itself)
- Claim: n will be expanded before B
 - $f(n)$ is less or equal to $f(A)$




$f(n) = g(n) + h(n)$	Definition of f-cost
$f(n) \leq g(A)$	Admissibility of h
$g(A) = f(A)$	$h = 0$ at a goal

Optimality of A* Tree Search

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, maybe A itself)
- Claim: n will be expanded before B
 - $f(n)$ is less or equal to $f(A)$
 - $f(A)$ is less than $f(B)$



$$g(A) < g(B)$$

B is suboptimal


$$f(A) < f(B)$$

$h = 0$ at a goal

Optimality of A* Tree Search

Proof:


- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 - $f(n)$ is less or equal to $f(A)$
 - $f(A)$ is less than $f(B)$
 - n expands before B
- All ancestors of A expand before B
- A expands before B
- A* search is optimal




$$f(n) \leq f(A) < f(B)$$

UCS vs A* Contours

- Uniform-cost expands equally in all “directions”



- A* expands mainly toward the goal, but does hedge its bets to ensure optimality



Graph Search

- Idea: never expand a state twice
- How to implement:
 - Tree search + set of expanded states (“closed set”)
 - Expand the search tree node-by-node, but...
 - Before expanding a node, check to make sure its state has never been expanded before
 - If not new, skip it, if new add to closed set
- Important: store the closed set as a set, not a list (performance)
- Can graph search wreck completeness? Why/why not?
- How about optimality?


Graph search

```

function Graph-Search(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← Insert (Make-Node(Initial-State[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← Remove-Front(fringe)
    if Goal-Test(problem, State[node]) then return node
    if State[node] is not in closed then
      add State[node] to closed
      fringe ← Insert All(Expand(node, problem), fringe)
  end

```

A* Graph Search Gone Wrong?



Consistency of Heuristics

- Main idea: estimated heuristic costs \leq actual costs
 - Admissibility: heuristic cost \leq actual cost to goal


$$h(A) \leq \text{actual cost from A to G} = h^*(A)$$
 - Consistency: heuristic “arc” cost \leq actual cost for each arc

$$h(A) - h(C) \leq \text{cost}(A \text{ to } C)$$
- Consequences of consistency:
 - The f value along a path never decreases $h(A) \leq \text{cost}(A \text{ to } C) + h(C)$
 - A* graph search is optimal

Optimality of A* Graph Search

Sketch of proof: consider what A* does with a consistent heuristic:

- Fact 1: In tree search, A* expands nodes in increasing total f value (f-contours)
- Fact 2: For every state s, nodes that reach s optimally are expanded before nodes that reach s suboptimally
- Result: A* graph search is optimal



Optimality

- Tree search:
 - A* is optimal if heuristic is admissible
 - UCS is a special case ($h = 0$)
- Graph search:
 - A* optimal if heuristic is consistent
 - UCS optimal ($h = 0$ is consistent)
- Consistency implies admissibility
- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems

A*: Summary

- A* uses both backward costs and (estimates of) forward costs
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems

A* Applications

- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...