

Outline


- Local Search and Optimization Problems
 - Hill-climbing
 - Simulated annealing
 - Genetic algorithms
- Local search in continuous spaces
 - Linear Programming

Iterative Improvement Search

- **For some problems**, the state description provides all the information required for a solution
 - path to the goal is not critical
 - path costs irrelevant
 - global maximum or minimum corresponds to the optimal solution
 - » local optimum may be “good enough”
- iterative improvement algorithms start with some configuration, and try modifications to improve the quality
 - 8-queens: number of un-attacked queens
 - Traveling Salesperson Problem (VLSI layout)
- analogy: state space as landscape with hills and valleys
 - “height” indicates value:
 - hill-climbing: always going uphill or hill descent
 - gradient ascent/descent

Local Search

- Tree search keeps unexplored alternatives on the fringe (ensures completeness)
- Local search: improve a single option until you can't make it better (no fringe!)
- New successor function: local changes




- Generally much faster and more memory efficient (but incomplete and suboptimal)

Local Search and Optimization Problems


- For some problem classes, it is sufficient to find a solution
 - the actual path to the solution is not relevant
- Memory requirements can be dramatically reduced
 - modifying the current state to advance the search
 - only information about the current state is kept
 - all information about previous states is discarded
 - » including paths to nodes in the search tree
 - may have serious consequences
 - » no path information kept => states may be re-visited (loops)
 - » impacts completeness, optimality
 - since only information about the current state is kept, such methods are called local

Hill Climbing Diagram



5

Hill Climbing Quiz (Maximize)



Part B
Write Pseudo Code

- Input: Optimization Problem
- $\text{Neighbors}(x)$ returns the set of neighbors of a state
 - $\text{HighVal}(y, z)$ returns the neighbor with the highest value
 - $\text{Val}(x)$ returns value of x

Starting from X, where do you end up ?

Starting from Y, where do you end up ?

Starting from Z, where do you end up ?

6

Hill-climbing (or gradient ascent/descent)

function **Hill-Climbing**(*problem*) returns a state that is a local maximum

inputs: *problem*, a problem

local variables: *current*, a node

neighbor, a node

```
current ← Initial-State [problem]  
loop do  
  neighbor ← a highest-valued successor of current  
  if Value[neighbor] ≤ Value[current] then return State[current]  
  current ← neighbor  
end
```

Problem: finds local maximum (or plateau)

How to ascend or descend the hill

Lots of variations: Examples

- All of the following – along with combinations

Variant	Definition
steepest-ascent	choose the highest-valued neighbor
stochastic choose	randomly from higher-valued neighbors
first-choice	choose the first higher-valued neighbor
random-restart	conduct hill climbing multiple times
local beam search	chooses the k highest-valued neighbors
--	
Simulated Annealing	
Tabu Search	

Focus: Overcoming locality

- Random restart – multiple random start states
- Allow suboptimal moves to broaden the search in some systematic way – Simulated Annealing
- Run multiple times but remember earlier start points and generate new start points – Tabu Search

Simulated Annealing

- Early on, higher "temperature": more likely to accept neighbors that are worse than current state
- Later on, lower "temperature": less likely to accept neighbors that are worse than current state

Simulated Annealing

function **SIMULATED ANNEALING**(problem, max):

 current = initial state of problem

 for t = 1 to max:

 T = TEMPERATURE(t)

 neighbor = random neighbor of current

ΔE = how much better neighbor is than current

 if $\Delta E > 0$:

 current = neighbor

 else with probability $e^{\Delta E/T}$ set current = neighbor

 return current

Simulated Annealing

- Idea: Escape local maxima by allowing downhill moves
 - But make them rarer as time goes on

```

function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to "temperature"
  local variables: current, a node
                    next, a node
                    T, a "temperature" controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 

```

11

memory structures form what is known as the tabu list, a set of rules and banned solutions used to filter which solutions will be admitted to the neighborhood

Tabu Search

- Similarly to Hill climbing Tabu search uses a local or neighborhood search procedure to iteratively move from one potential solution to an improved solution in the neighborhood of x until some stopping criterion has been satisfied (generally, an attempt limit or a score threshold).
- To avoid this and explore other regions of the search space, tabu search carefully explores the neighborhood of each solution as the search progresses. The solutions admitted to the new neighborhood, $N^*(x)$, are determined using memory structures.
- Using these memory structures, the search progresses by iteratively moving from the current solution x to an improved solution x' in $N^*(x)$

12

Tabu Search (continued)


The memory structures used in tabu search can roughly be divided into three categories:

- Short-term: The list of solutions recently considered. If a potential solution appears on the tabu list, it cannot be revisited until it reaches an **expiration time**.
- Intermediate-term: **Intensification rules** are intended to bias the search towards promising areas of the search space.
- Long-term: **Diversification rules** that drive the search into new regions (i.e., regarding resets when the search becomes stuck in a plateau or a suboptimal dead-end).

13

Example: n-queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
- Move a queen to reduce number of conflicts




- Almost always solves n-queens problems almost instantaneously for very large n , e.g., $n = 1$ million

14


Traveling Salesman Problem

- Given n cities with known distances between each pair, find the shortest tour that passes through all the cities exactly once before returning to the starting city

An Instance of the Traveling Salesman Problem





Search Space



15

Local Search Heuristics for TSP

- Idea: Iteratively improve on an approximation algorithm
 - Start with some initial tour (e.g., nearest neighbor). On each iteration, explore the current tour's neighborhood by exchanging a few edges in it. If the new tour is shorter, make it the current tour; otherwise consider another edge change. If no change yields a shorter tour, the current tour is returned as the output
- Two Opt



16

Local beam search

- Idea: keep k states instead of 1; choose top k of all their successors
- Not the same as k searches run in parallel!
- Searches that find good states recruit other searches to join them
- Problem: quite often, all k states end up on same local hill Idea: choose k successors randomly, biased towards good ones
Observe the close analogy to natural selection!

Continuous State Spaces- revisit when get to neural networks

Suppose we want to site three airports in Romania:

- 6-D state space defined by $(x_1, y_1), (x_2, y_2), (x_3, y_3)$
- objective function $f(x_1, y_1, x_2, y_2, x_3, y_3) =$
sum of squared distances from each city to nearest airport

Gradient methods compute

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial y_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial y_2}, \frac{\partial f}{\partial x_3}, \frac{\partial f}{\partial y_3} \right)$$

to increase/reduce f , e.g., by $x \leftarrow x + a \nabla f(x)$

Sometimes can solve for $\nabla f(x) = 0$ exactly (e.g., with one city).

Newton-Raphson (1664, 1690) iterates $x \leftarrow x - H_f^{-1}(x) \nabla f(x)$
to solve $\nabla f(x) = 0$, where $H_{ij} = \partial^2 f / \partial x_i \partial x_j$

Linear Programming (LP) problem

Optimize a linear function of several variables subject to linear constraints:

Maximize (or Minimize)

$$c_1 x_1 + \dots + c_n x_n$$

subject to

$$a_{i,1} x_1 + \dots + a_{i,n} x_n \leq b_i, i=1, \dots, m$$

$$x_1 \geq 0, \dots, x_n \geq 0$$


- The function $z = c_1 x_1 + \dots + c_n x_n$ is called the objective function; x_i : decision variables
- The constraints $x_1 \geq 0, \dots, x_n \geq 0$ are called non-negativity constraints

19

Example: Want to maximize the profit on two products where production is constrained.

$$\begin{array}{ll} \text{maximize} & 3x + 5y \\ \text{subject to} & x + y \leq 4 \\ & x + 3y \leq 6 \\ & x \geq 0, y \geq 0 \end{array}$$

Feasible region is the set of points defined by the constraints




20

Geometric solution

maximize $3x + 5y$
 subject to $x + y \leq 4$
 $x + 3y \leq 6$
 $x \geq 0, y \geq 0$

Optimal solution: $x = 3, y = 1$



Extreme Point Theorem: Any LP problem with a nonempty bounded feasible region has an optimal solution; moreover, an optimal solution can always be found at an extreme point of the problem's feasible region.

3 possible outcomes in solving an LP problem



- has a finite optimal solution, which may not be unique
- unbounded: the objective function of maximization (minimization) LP problem is unbounded from above (below) on its feasible region
- infeasible: there are no points satisfying all the constraints, i.e. the constraints are contradictory

The Chocolate Shop

- You own a chocolate shop which produces two types of box chocolates:
 - » Normal box which gives a \$1 profit
 - » Deluxe box which gives a \$6 profit
- The variables are the number of boxes produced per day
 - » x_1 is the number of boxes of normal chocolate
 - » x_2 is the number of boxes of deluxe chocolate
- The objective is to set x_1 and x_2 to maximize profit
 - » $\max (x_1 + 6x_2)$ Profit = $x_1 + 6x_2$
- The constraints are:
 - » $x_1 \leq 200$ Maximum demand of normal boxes per day
 - » $x_2 \leq 300$ Maximum demand of deluxe boxes per day
 - » $x_1 + x_2 \leq 400$ Maximum production capacity
 - » $x_1, x_2 \geq 0$ Can't have a negative number of boxes


Representing a Linear Program

- The objectives and constraints define a linear program
- Easy to visualize in low dimensions (2 or 3 variables)
 - Feasible space forms a convex polygon
- Optimum is achieved at a vertex, except when
 - No solution to the constraints
 - Feasible region is unbounded in direction of the objective



Solving a Linear Program

- Simplex Algorithm
 - Start at any vertex
 - If a neighbor has a better objective value move to it, otherwise already at optimum
 - » Objective value = $x_1 + 6x_2$
- Is this optimal?
 - Convexity and linearity the reason
- If more than one neighbor with a better objective value, which one should you go to?



Slides and Course Materials modified from

University of California, Berkeley [These slides adapted from Dan Klein and Pieter Abbeel]

CS50 Course from Harvard


???

CSC 480 from Cal Poly

Franz Kurfess

Genetic algorithms

GA : stochastic local beam search +
generate successors from pairs of states



Genetic algorithms: Example

- GAs require states encoded as strings (GPs use programs)
- Crossover helps if and only if substrings are meaningful components
- GAs \neq evolution: e.g., real genes encode replication machinery!

